

# UCORE Porting on MIPS32S CPU

Chen Yuheng  
Tsinghua Univ.

November 26, 2013

## 1 Introduction

This report gives a brief introduction to our project – *UCORE* porting onto the MIPS32S platform, including the kernel, limited device drivers as well as the user-lib. *UCORE* is an experimental modern operating system that includes a basic memory manager, a limited process/ thread scheduler, linux-like VFS and incomplete POSIX-compatible syscall interface. Currently, the following platforms are supported:

- *Mipssim in Qemu*, which is simulated with qemu-system-mipsel. The hardware configuration is available in qemu’s source code, several important components of which is summarized in Table. 1.

Component	Type	Base Address	IRQ
Uart	16450	0x1fd003f8	4
Timer	On Chip	–	7
PIC	On Chip	–	–

Table 1: Mipssim Hardware Configuration

- *Our FPGA MIPS32S Implementation* hardware platform.

Our code is available on Github <sup>1</sup> and is kept updating. Any bug reports are welcomed.

## 2 Architecture of the MIPS32S Platform

## 3 Development Environment

This section is a guide to setup and use the cross-compile development environment for Mips.

---

<sup>1</sup><https://github.com/chyh1990/ucore-thumips>

### 3.1 Toolchain

We use the standard toolchain GCC for Mips to compile ucore, even on our MIPS32S CPU, in which only a subset of MIPS1 Instruction Set is implemented. The best place to get a free GCC-MIPS toolchain is on CodeSourcery <sup>2</sup>, just make sure you download GCC 4.6. The package also includes gdb for MIPS.

Another way to get a toolchain is compiling it from source code. GCC-core 4.6.3+Binutils 2.22 is tested. Compiling GCC is tricky, just google it if you come into any issues. Here is my configuration for GCC on 64-bit Ubuntu 12.04:

```
../gcc-4.6.3/configure --prefix=/home/guest/cpu/build-gcc/mips_gcc
--target=mips-sde-elf --disable-nls --enable-languages=c
--disable-multilib --without-headers --disable-shared
--without-newlib --disable-libgomp --disable-libssp
--disable-threads --disable-libgcc
```

### 3.2 Simulator

You can run ucore-thumips with the standard Qemu in your Linux distribution:

```
qemu-system-mipsel -M mipssim -m 32M -serial stdio
-kernel obj/ucore-kernel-initrd
```

But we recommend to use our modified Qemu to simulate our simplified MIPS Instruction Set and Flash support.

what's more, it will be necessary to setup the mips-sde-elf-gdb properly, refer to the following *.gdbinit* example:

```
set endian little
set mipsfpu none
target remote 127.0.0.1:1234
file obj/ucore-kernel-initrd
```

### 3.3 MIPS32S Programming Guide

MIPS32S supports a simplified MIPS32 Instruction Set, here is several important differences:

1. No *lh/sh*
2. No *divu*
3. No *add/sub*, only *addu/subu* is supported.

We use some tricks to solve these problems, see *kern/thumips.h*.

Another problem is that MIPS32S does NOT support delayed slot. Solving this problem by giving GCC options correctly:

```
CFLAGS := -EL -G0 -fno-delayed-branch -Wa,-O0
```

An example MIPS32S C project can be found in test1.tar.gz.

---

<sup>2</sup><https://sourcery.mentor.com/GNUToolchain/release2189>

## 4 Source Code Organization

This section introduces the source code organization of ucore-thumpis and explains several configuration options.

### 4.1 Source Tree

Since our work is based on LAB0-LAB8, important directories are listed in Table. 2.

Directory	Description
debug	debug console after a kernel panic
driver	device driver interface definition
include	useful macros for MIPS32S
fs	Filesystem
init	kernel entry point and initialization code
libs	utilities
mm	low-level memory management
process	context switch
sync	atomic operation
syscall	MIPS-specific syscall mechanism
trap	exception handling

Table 2: ucore-thumpis directories

### 4.2 Makefile

UCore does not have a configuration system yet, all configuration is hand-coded in the Makefile.

1. *GCCPREFIX*, toolchain path;
2. *USER\_APPLIST*, the applications to be included.

More detailed memory layout is defined in memlayout.h in the corresponding machine directory, see Section. 5.3.

## 5 Implementation Details

This section describes the implementations of some important mechanism in the MIPS32S CPU, which is similar to the standard MIPS32 CPU. Thus, I refer to Harvard's OS/161[1] project during writing the code.

### 5.1 Booting

The booting process is also simulated. Our Qemu needs the following files: There are two ways to use our modified Qemu to boot ucoer for MIPS. Then, the C environment must be setup:

- *sp*, setup the stack

File	Type	Description
ucore-kernel-initrd	ELF	ucore kernel with ramdisk
flash.img	Binary	used for flash simulation, can be a link to the kernel
boot/loader.bin	Binary	BIOS, knows ELF and load kernel from Flash
thumips_insn.txt	Text	Instruction Set Descriptor

Table 3: Qemu Files

Offset	Size	Usage
0x00000000	0x80000000	Userspace, TLB mapped
0x80000000	0x20000000	Kernel Space, Direct mapping
0xA0000000	0x20000000	IO Space, Direct Mapping

Table 4: Bootable Kernel Memory Layout

- *gp*, set to *\_gp* (defined in *ldscript*)
- zero *.bss* section

For convenience, root filesystem image(ramdisk) is linked with the kernel, appending at the end of the *.data* section. <sup>3</sup>.

## 5.2 Exception Handling

The most important hardware support for exception/interrupt handling on MIPS32S is the *SR* register:




Figure 1: Status Register[2]

- ERL, MUST be clear by software after reset;
- EXL and KSU, MUST be set properly to handle nested exceptions.

The kernel use a trapframe structure to save registers when privilege mode transition occurs:

The trapframe is constructed by the assembly in *trap/exception.S*, be aware with the fact that registers *k0* and *k1* is reserved by the compiler so we can use them in the exception handler without saving them.

See *kern/trap/exception.S* for more details.

## 5.3 Memory Management

MIPS32S has no MMUs, but has a programmable TLB unit. So it is possible to emulate MMU utilizing the TLB miss exception. The software simulated

<sup>3</sup>see *tools/kernel.ld*

---

**Algorithm 1** Trapframe

---

```
/* $1 - $30 */
struct pushregs {
    uint32_t reg_r[30];
};

/*
 * Structure describing what is saved on the stack during entry to
 * the exception handler.
 *
 * This must agree with the code in exception.S.
 */

struct trapframe {
    uint32_t tf_vaddr; /* coprocessor 0 vaddr register */
    uint32_t tf_status; /* coprocessor 0 status register */
    uint32_t tf_cause; /* coprocessor 0 cause register */
    uint32_t tf_lo;
    uint32_t tf_hi;
    uint32_t tf_ra; /* Saved register 31 */
    struct pushregs tf_regs;
    uint32_t tf_epc; /* coprocessor 0 epc register */
};
```

---

MMU of MIPS32S is similar to X86's, Ucore uses the following configuration (which is similar to X86's VA layout) :

PDT Index	PTE Index	Offset
10	10	12

Table 5: Virtual Address in ucore

In our TLB miss handler, we first checkout whether it is just a TLB miss or a page table miss by looking the address up in our emulated X86 page table. If it is a TLB miss, check the access permission and fill it up (or raise a access violation). If it is a page miss, call the do\_pgfault. See Alg. 2.

## 5.4 Context Switch

According to MIPS32 O32 ABI, we must save the following registers during the context switching:

---

**Algorithm 3** Context

---

```
struct context {  
    uint32_t sf_s0;  
    uint32_t sf_s1;  
    uint32_t sf_s2;  
    uint32_t sf_s3;  
    uint32_t sf_s4;  
    uint32_t sf_s5;  
    uint32_t sf_s6;  
    uint32_t sf_s7;  
    uint32_t sf_s8;  
    uint32_t sf_gp;  
    uint32_t sf_ra;  
    uint32_t sf_sp;  
};
```

---

According to O32 ABI, *s0-s8* must be reserved by the callee, *gp* is a global pointer, *ra* is the interrupted PC. In addition, we must switch the kernel stack, which is saved in *sf\_sp*.

## 5.5 System Call

The syscall mechanism is borrowed from Linux2.6. In MIPS, a special instruction *syscall* handles user mode to supervisor mode transition.

In ucore, we employ the following system calling convention:

1. *a0 - a3*, arguments(from left to right)
2. *v0*, syscall number
3. *syscall*
4. return value in *v0*

Syscall is wrapped in user mode library and can be called as a normal C function.

## 5.6 User Library and Application

The code in *user* should works without modification. However, since *libs-user-ucore/syscall.c* is not compatible with MIPS's system calling convention, we use some macros from Linux to create our own system call entries. See Alg. 4.

Another modification is *user/libs/user.ld*. The .text section of user application is located at 0x10000000.

## 6 Conclusion

In our project, we work out a basically working version of ucore for MIPS32S. However, just as what items listed on the TODO list imply, much work remains for making ucore for MIPS32S practically usable operating system.

## References

- [1] David A. Holland, Ada T. Lim, and Margo I. Seltzer. A new instructional operating system. *SIGCSE Bull.*, 34(1):111–115, February 2002.
- [2] Dominic Sweetman. *See MIPS Run, Second Edition*. Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 2006.

---

**Algorithm 2** User Mode System Calling Convetion

---

```
/* use software emulated X86 pgfault */
static void handle_tlbmiss(struct trapframe* tf, int write)
{
    int in_kernel = trap_in_kernel(tf);
    assert(current_pgdir != NULL);
    uint32_t badaddr = tf->tf_vaddr;
    int ret = 0;
    pte_t *pte = get_pte(current_pgdir, tf->tf_vaddr, 0);
    if(pte==NULL || ptep_invalid(pte)){ //PTE miss, pgfault
        //tlb will not be refill in do_pgfault,
        //so a vmm pgfault will trigger 2 exception
        //permission check in tlb miss
        ret = pgfault_handler(tf, badaddr, get_error_code(write, pte));
    }else{ //tlb miss only, reload it
        /* refill two slot */
        /* check permission */
        if(in_kernel){
            tlb_refill(badaddr, pte);
            return;
        }else{
            if(!ptep_u_read(pte)){
                ret = -1;
                goto exit;
            }
            if(write && !ptep_u_write(pte)){
                ret = -2;
                goto exit;
            }
            tlb_refill(badaddr, pte);
            return ;
        }
    }
}

exit:
    if(ret){
        print_trapframe(tf);
        if(in_kernel){
            panic("unhandled pgfault");
        }else{
            do_exit(-E.KILLED);
        }
    }
    return ;
}
```

---



---

**Algorithm 4** User Mode System Calling Convnetion

---

```
static inline int
syscall(int num, ...) {
    va_list ap;
    va_start(ap, num);
    uint32_t arg[MAX_ARGS];
    int i, ret;
    for (i = 0; i < MAX_ARGS; i++) {
        arg[i] = va_arg(ap, uint32_t);
    }
    va_end(ap);

    num += SYSCALL_BASE;

    asm volatile(
        ".set noreorder;\n"
        "move $v0, %1;\n" /* syscall no. */
        "move $a0, %2;\n"
        "move $a1, %3;\n"
        "move $a2, %4;\n"
        "move $a3, %5;\n"
        "syscall;\n"
        "nop;\n"
        "move %0, $v0;\n"
        : "=r"(ret)
        : "r"(num), "r"(arg[0]), "r"(arg[1]), "r"(arg[2]), "r"(arg
          [3])
        : "a0", "a1", "a2", "a3", "v0"
    );
    return ret;
}
```

---